

open(3) - Linux man page

Prolog

This manual page is part of the POSIX Programmer's Manual. The Linux implementation of this interface may differ (consult the corresponding Linux manual page for details of Linux behavior), or the interface may not be implemented on Linux.

Name

open - open a file

Synopsis

#include <[sys/stat.h](#)>

#include <[fcntl.h](#)>

```
int open(const char *path, int oflag, ... );
```

Description

The *open()* function shall establish the connection between a file and a file descriptor. It shall create an open file description that refers to a file and a file descriptor that refers to that open file description. The file descriptor is used by other I/O functions to refer to that file. The *path* argument points to a pathname naming the file.

The *open()* function shall return a file descriptor for the named file that is the lowest file descriptor not currently open for that process. The open file description is new, and therefore the file descriptor shall not share it with any other process in the system. The `FD_CLOEXEC` file descriptor flag associated with the new file descriptor shall be cleared.

The file offset used to mark the current position within the file shall be set to the beginning of the file.

The file status flags and file access modes of the open file description shall be set according to the value of *oflag*.

Values for *oflag* are constructed by a bitwise-inclusive OR of flags from the following list, defined in <[fcntl.h](#)>. Applications shall specify exactly one of the first three values (file access modes) below in the value of *oflag*:

O_RDONLY

Open for reading only.

O_WRONLY

Open for writing only.

O_RDWR

Open for reading and writing. The result is undefined if this flag is applied to a FIFO.

Any combination of the following may be used:

O_APPEND

If set, the file offset shall be set to the end of the file prior to each write.

O_CREAT

If the file exists, this flag has no effect except as noted under O_EXCL below. Otherwise, the file shall be created; the user ID of the file shall be set to the effective user ID of the process; the group ID of the file shall be set to the group ID of the file's parent directory or to the effective group ID of the process; and the access permission bits (see [*<sys/stat.h>*](#)) of the file mode shall be set to the value of the third argument taken as type **mode_t** modified as follows: a bitwise AND is performed on the file-mode bits and the corresponding bits in the complement of the process' file mode creation mask. Thus, all bits in the file mode whose corresponding bit in the file mode creation mask is set are cleared. When bits other than the file permission bits are set, the effect is unspecified. The third argument does not affect whether the file is open for reading, writing, or for both. Implementations shall provide a way to initialize the file's group ID to the group ID of the parent directory. Implementations may, but need not, provide an implementation-defined way to initialize the file's group ID to the effective group ID of the calling process.

O_DSYNC

Write I/O operations on the file descriptor shall complete as defined by synchronized I/O data integrity completion.

O_EXCL

If O_CREAT and O_EXCL are set, *open()* shall fail if the file exists. The check for the existence of the file and the creation of the file if it does not exist shall be atomic with respect to other threads executing *open()* naming the same filename in the same directory with O_EXCL and O_CREAT set. If O_EXCL and O_CREAT are set, and *path* names a symbolic link, *open()* shall fail and set *errno* to [EEXIST], regardless of the contents of the symbolic link. If O_EXCL is set and O_CREAT is not set, the result is undefined.

O_NOCTTY

If set and *path* identifies a terminal device, *open()* shall not cause the terminal device to become the controlling terminal for the process.

O_NONBLOCK

When opening a FIFO with O_RDONLY or O_WRONLY set:

*

If O_NONBLOCK is set, an *open()* for reading-only shall return without delay. An *open()* for writing-only shall return an error if no process currently has the file open for reading.

*

If O_NONBLOCK is clear, an *open()* for reading-only shall block the calling thread until a thread opens the file for writing. An *open()* for writing-only shall block the calling thread until a thread opens the file for reading.

When opening a block special or character special file that supports non-blocking opens:

*

If O_NONBLOCK is set, the *open()* function shall return without blocking for the device to be ready or available. Subsequent behavior of the device is device-specific.

*

If O_NONBLOCK is clear, the *open()* function shall block the calling thread until the device is ready or available before returning.

Otherwise, the behavior of O_NONBLOCK is unspecified.

O_RSYNC

Read I/O operations on the file descriptor shall complete at the same level of integrity as specified by the O_DSYNC and O_SYNC flags. If both O_DSYNC and O_RSYNC are set in *oflag*, all I/O operations on the file descriptor shall complete as defined by synchronized I/O data integrity completion. If both O_SYNC and O_RSYNC are set in flags, all I/O operations on the file descriptor shall complete as defined by synchronized I/O file integrity completion.

O_SYNC

Write I/O operations on the file descriptor shall complete as defined by synchronized I/O file integrity completion.

O_TRUNC

If the file exists and is a regular file, and the file is successfully opened O_RDWR or O_WRONLY, its length shall be truncated to 0, and the mode and owner shall be unchanged. It shall have no effect on FIFO special files or

terminal device files. Its effect on other file types is implementation-defined. The result of using `O_TRUNC` with `O_RDONLY` is undefined.

If `O_CREAT` is set and the file did not previously exist, upon successful completion, `open()` shall mark for update the `st_atime`, `st_ctime`, and `st_mtime` fields of the file and the `st_ctime` and `st_mtime` fields of the parent directory.

If `O_TRUNC` is set and the file did previously exist, upon successful completion, `open()` shall mark for update the `st_ctime` and `st_mtime` fields of the file.

If both the `O_SYNC` and `O_DSYNC` flags are set, the effect is as if only the `O_SYNC` flag was set.

If *path* refers to a STREAMS file, *oflag* may be constructed from `O_NONBLOCK` OR'ed with either `O_RDONLY`, `O_WRONLY`, or `O_RDWR`. Other flag values are not applicable to STREAMS devices and shall have no effect on them. The value `O_NONBLOCK` affects the operation of STREAMS drivers and certain functions applied to file descriptors associated with STREAMS files. For STREAMS drivers, the implementation of `O_NONBLOCK` is device-specific.

If *path* names the master side of a pseudo-terminal device, then it is unspecified whether `open()` locks the slave side so that it cannot be opened. Conforming applications shall call `unlockpt()` before opening the slave side.

The largest value that can be represented correctly in an object of type **off_t** shall be established as the offset maximum in the open file description.

Return Value

Upon successful completion, the function shall open the file and return a non-negative integer representing the lowest numbered unused file descriptor. Otherwise, -1 shall be returned and *errno* set to indicate the error. No files shall be created or modified if the function returns -1.

Errors

The `open()` function shall fail if:

EACCES

Search permission is denied on a component of the path prefix, or the file exists and the permissions specified by *oflag* are denied, or the file does not exist and write permission is denied for the parent directory of the file to be created, or `O_TRUNC` is specified and write permission is denied.

EEXIST

O_CREAT and O_EXCL are set, and the named file exists.

EINTR

A signal was caught during *open()*.

EINVAL

The implementation does not support synchronized I/O for this file.

EIO

The *path* argument names a STREAMS file and a hangup or error occurred during the *open()*.

EISDIR

The named file is a directory and *oflag* includes O_WRONLY or O_RDWR.

ELOOP

A loop exists in symbolic links encountered during resolution of the *path* argument.

EMFILE

{OPEN_MAX} file descriptors are currently open in the calling process.

ENAMETOOLONG

The length of the *path* argument exceeds {PATH_MAX} or a pathname component is longer than {NAME_MAX}.

ENFILE

The maximum allowable number of files is currently open in the system.

ENOENT

O_CREAT is not set and the named file does not exist; or O_CREAT is set and either the path prefix does not exist or the *path* argument points to an empty string.

ENOSR

The *path* argument names a STREAMS-based file and the system is unable to allocate a STREAM.

ENOSPC

The directory or file system that would contain the new file cannot be expanded, the file does not exist, and O_CREAT is specified.

ENOTDIR

A component of the path prefix is not a directory.

ENXIO

O_NONBLOCK is set, the named file is a FIFO, O_WRONLY is set, and no process has the file open for reading.

ENXIO

The named file is a character special or block special file, and the device associated with this special file does not exist.

EOVERFLOW

The named file is a regular file and the size of the file cannot be represented correctly in an object of type **off_t**.

EROFS

The named file resides on a read-only file system and either `O_WRONLY`, `O_RDWR`, `O_CREAT` (if the file does not exist), or `O_TRUNC` is set in the *oflag* argument.

The *open()* function may fail if:

EAGAIN

The *path* argument names the slave side of a pseudo-terminal device that is locked.

EINVAL

The value of the *oflag* argument is not valid.

ELOOP

More than `{SYMLOOP_MAX}` symbolic links were encountered during resolution of the *path* argument.

ENAMETOOLONG

As a result of encountering a symbolic link in resolution of the *path* argument, the length of the substituted pathname string exceeded `{PATH_MAX}`.

ENOMEM

The *path* argument names a STREAMS file and the system is unable to allocate resources.

ETXTBSY

The file is a pure procedure (shared text) file that is being executed and *oflag* is `O_WRONLY` or `O_RDWR`.

The following sections are informative.

Examples

Opening a File for Writing by the Owner

The following example opens the file **/tmp/file**, either by creating it (if it does not already exist), or by truncating its length to 0 (if it does exist). In the former case, if the call creates a new file, the access permission bits in the file mode of the file are set to permit reading and writing by the owner, and to permit reading only by group members and others.

If the call to *open()* is successful, the file is opened for writing.

```

#include <fcntl.h>
...
int fd;
mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
char *filename = "/tmp/file";
...
fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, mode);
...

```

Opening a File Using an Existence Check

The following example uses the *open()* function to try to create the **LOCKFILE** file and open it for writing. Since the *open()* function specifies the *O_EXCL* flag, the call fails if the file already exists. In that case, the program assumes that someone else is updating the password file and exits.

```

#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>

#define LOCKFILE "/etc/ptmp"
...
int pfd; /* Integer for file descriptor returned by open() call. */
...
if ((pfd = open(LOCKFILE, O_WRONLY | O_CREAT | O_EXCL,
    S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)) == -1)
{
    fprintf(stderr, "Cannot open /etc/ptmp. Try again later.\n");
    exit(1);
}
...

```

Opening a File for Writing

The following example opens a file for writing, creating the file if it does not already exist. If the file does exist, the system truncates the file to zero bytes.

```

#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>

```

```

#define LOCKFILE "/etc/ptmp"
...
int pfd;
char filename[PATH_MAX+1];
...
if ((pfd = open(filename, O_WRONLY | O_CREAT | O_TRUNC,
    S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)) == -1)
{
    perror("Cannot open output file\n"); exit(1);
}
...

```

Application Usage

None.

Rationale

Except as specified in this volume of IEEE Std 1003.1-2001, the flags allowed in *oflag* are not mutually-exclusive and any number of them may be used simultaneously.

Some implementations permit opening FIFOs with O_RDWR. Since FIFOs could be implemented in other ways, and since two file descriptors can be used to the same effect, this possibility is left as undefined.

See *getgroups()* about the group of a newly created file.

The use of *open()* to create a regular file is preferable to the use of *creat()*, because the latter is redundant and included only for historical reasons.

The use of the O_TRUNC flag on FIFOs and directories (pipes cannot be *open()*-ed) must be permissible without unexpected side effects (for example, *creat()* on a FIFO must not remove data). Since terminal special files might have type-ahead data stored in the buffer, O_TRUNC should not affect their content, particularly if a program that normally opens a regular file should open the current controlling terminal instead. Other file types, particularly implementation-defined ones, are left implementation-defined.

IEEE Std 1003.1-2001 permits [EACCES] to be returned for conditions other than those explicitly listed.

The `O_NOCTTY` flag was added to allow applications to avoid unintentionally acquiring a controlling terminal as a side effect of opening a terminal file. This volume of IEEE Std 1003.1-2001 does not specify how a controlling terminal is acquired, but it allows an implementation to provide this on *open()* if the `O_NOCTTY` flag is not set and other conditions specified in the Base Definitions volume of IEEE Std 1003.1-2001, Chapter 11, General Terminal Interface are met. The `O_NOCTTY` flag is an effective no-op if the file being opened is not a terminal device.

In historical implementations the value of `O_RDONLY` is zero. Because of that, it is not possible to detect the presence of `O_RDONLY` and another option. Future implementations should encode `O_RDONLY` and `O_WRONLY` as bit flags so that:

$$\text{O_RDONLY} \mid \text{O_WRONLY} == \text{O_RDWR}$$

In general, the *open()* function follows the symbolic link if *path* names a symbolic link. However, the *open()* function, when called with `O_CREAT` and `O_EXCL`, is required to fail with `[EEXIST]` if *path* names an existing symbolic link, even if the symbolic link refers to a nonexistent file. This behavior is required so that privileged applications can create a new file in a known location without the possibility that a symbolic link might cause the file to be created in a different location.

For example, a privileged application that must create a file with a predictable name in a user-writable directory, such as the user's home directory, could be compromised if the user creates a symbolic link with that name that refers to a nonexistent file in a system directory. If the user can influence the contents of a file, the user could compromise the system by creating a new system configuration or spool file that would then be interpreted by the system. The test for a symbolic link which refers to a nonexistent file must be atomic with the creation of a new file.

The POSIX.1-1990 standard required that the group ID of a newly created file be set to the group ID of its parent directory or to the effective group ID of the creating process. FIPS 151-2 required that implementations provide a way to have the group ID be set to the group ID of the containing directory, but did not prohibit implementations also supporting a way to set the group ID to the effective group ID of the creating process. Conforming applications should not assume which group ID will be used. If it matters, an application can use *chown()* to set the group ID after the file is created, or determine under what conditions the implementation will set the desired group ID.

Future Directions

[Translate](#)

None.

See Also

chmod(), *close()*, *creat()*, *dup()*, *fcntl()*, *lseek()*, *read()*, *umask()*, *unlockpt()*, *write()*, the Base Definitions volume of IEEE Std 1003.1-2001, [*<fcntl.h>*](#), [*<sys/stat.h>*](#), [*<sys/types.h>*](#)

Copyright

Portions of this text are reprinted and reproduced in electronic form from IEEE Std 1003.1, 2003 Edition, Standard for Information Technology -- Portable Operating System Interface (POSIX), The Open Group Base Specifications Issue 6, Copyright 漏 2001-2003 by the Institute of Electrical and Electronics Engineers, Inc and The Open Group. In the event of any discrepancy between this version and the original IEEE and The Open Group Standard, the original IEEE and The Open Group Standard is the referee document. The original Standard can be obtained online at <http://www.opengroup.org/unix/online.html> .

Referenced By

[**shortrpm**](#)(1)