

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
```

```
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE * f;
```

```
    for(i=0; i<MAXPAROLA; i++)
        freq[i]=0;
```

```
    if(argc != 2)
```

```
    {
        fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
        exit(1);
    }
```

```
    f = fopen(argv[1], "r");
    if(f==NULL)
```

```
    {
        fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
        exit(1);
    }
```

```
    while( fgets( riga, MAXRIGA, f ) != NULL )
```

Processes

Processes and concurrency

Stefano Quer and Pietro Laface

Dipartimento di Automatica e Informatica

Politecnico di Torino

Program

- ❖ **Algorithm:** a logical procedure that in a finite number of steps solves a problem
- ❖ **Program:** formal expression of an algorithm by means of a programming language
 - Static entity
 - Sequence of code lines

Process

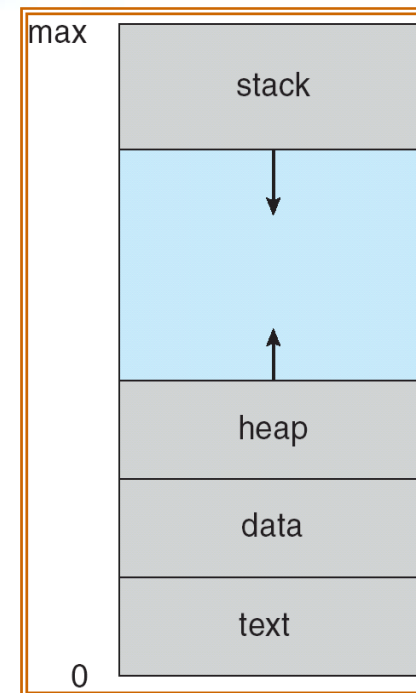
- ❖ **Process**: a sequence of operations performed by a program in execution on a given set of input data.
- ❖ The temporal behavior of a process can be analyzed through its **trace**
 - Program in execution
 - Dynamic entity

trace execute like debug

Program counter	Stack pointer	Register A	Register B	Variable X	Variable Y
0	0x1234	0	0	0	0
4	0x1234	-10	0	0	0
8	0x1234	-10	0	-10	0

Process

- Text area (source code) **executable code**
- Data area (global variables)
- Stack (function parameters and local variables)
- Heap (dynamic variables allocated during the process execution)
- Registers (Program counter, stack pointer, etc.)

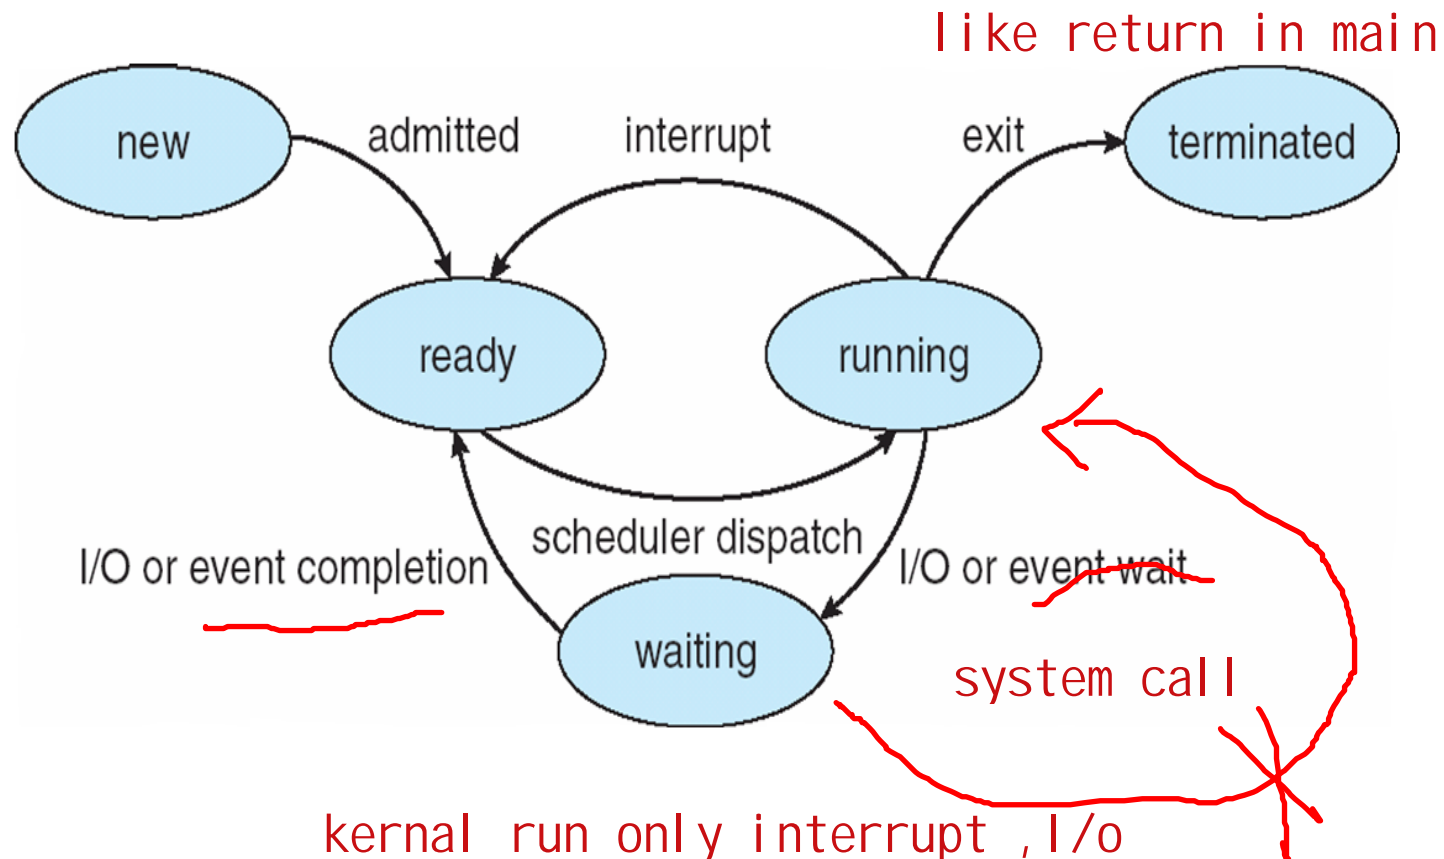
**local var****malloc****total memory**

Process state

- ❖ During its execution a process change its state
 - **New:** process is created and submitted to the OS ask for res
 - **Running:** a CPU is allocated to the process
 - **Ready:** logically ready to run, waiting that a CPU is available
 - **Waiting:** for an event or for a resource
 - **Terminated:** releases the resource it is using

State diagram

- ❖ The possible state evolution of a process is described by a state diagram



Process Control Block (PCB)

- ❖ The kernel stores for each process a set of data, e.g.,
- The process state
 - New, Ready, Running, Waiting, Terminated
 - Copy of the CPU registers
 - Their number and type is hardware-dependent
 - The program counter
 - Address of the next instruction to be executed

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

where it can access

snapshot of hardware situation of trace
for restart the process

Process Control Block (PCB)

- Data useful for CPU scheduling
 - Priority, pointers to queues, etc.
- Data useful for memory management
 - Segment and paging registers, segment and page tables, etc.
- File table
 - open files
- Signal table
 - signal handlers
- Etc.

link to some state list

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

Process Control Block (PCB)

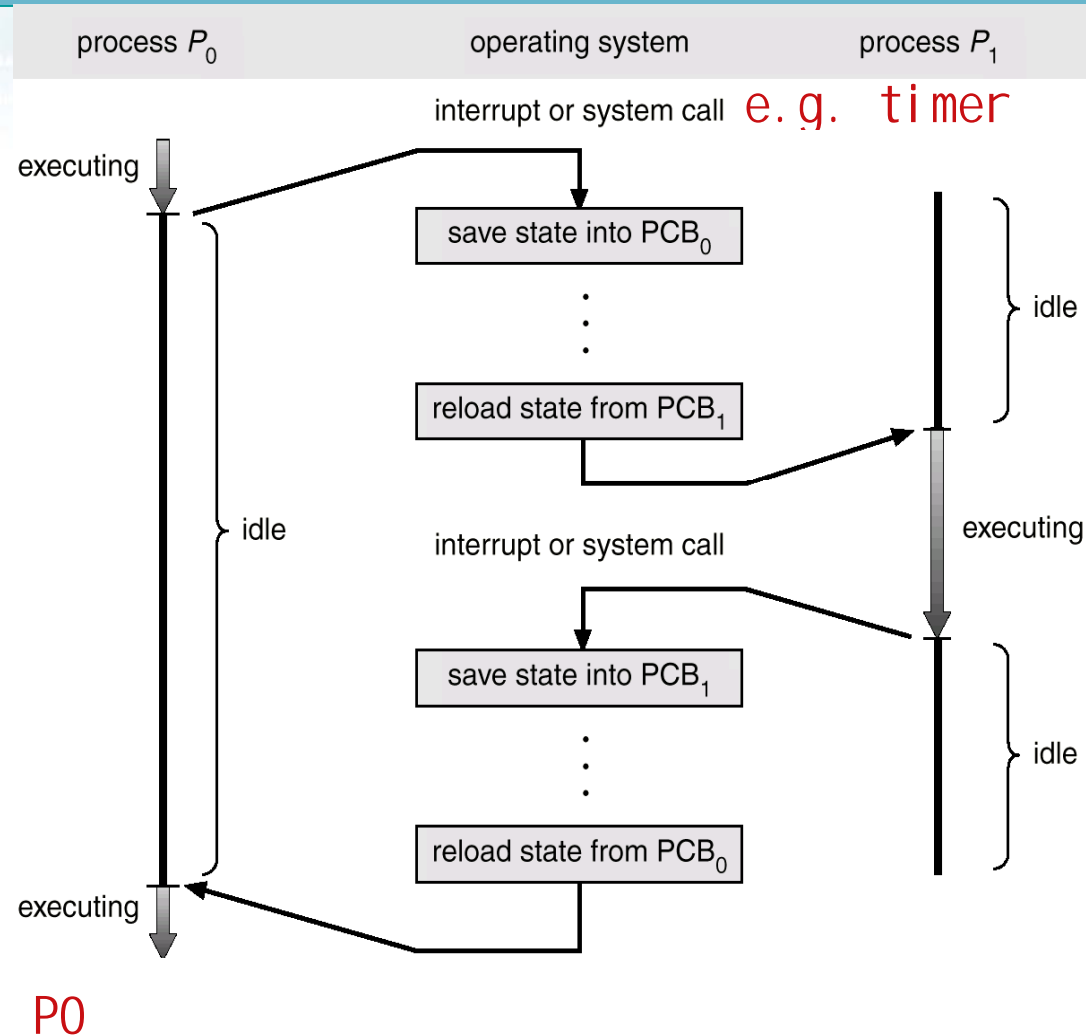
- Various administration data
 - CPU usage, limits, etc.
- I/O status information
 - I/O device list, etc.

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

Context switching

- ❖ When the CPU is assigned to another process, the kernel
 - Save the state of the running process **snapshot**
 - Load the state previously saved for the new process
- ❖ The time devoted to this **context switching** is overhead, i.e., time not directly useful for any process
- ❖ The amount of time for context switching is hardware-dependent **application independent**

Context switching



Process scheduling

- ❖ Multiprogramming aims at maximizing the CPU usage by processes
- ❖ Processes can be classified as
 - I/O-bound
 - Spend more time for I/O than for computation
 - Require short CPU service times
 - CPU-bound
 - Spend more time for computation than for I/O
 - Require long CPU service times

Process scheduling

- ❖ The kernel manages the sharing of the CPU among processes by means of a **scheduler**
 - A scheduler **selects the next process to run**, among the ready ones, according to a strategy that tries to maximize the CPU usage and to satisfy the response time for users

Process scheduling

❖ Different types of schedulers

➤ Short-term scheduler

- Selects the next process to run (context-switching)
- Run frequently
- Rescheduling performed every 1 to 10 milliseconds
- Must be extremely fast

Process scheduling

➤ Long-term scheduler

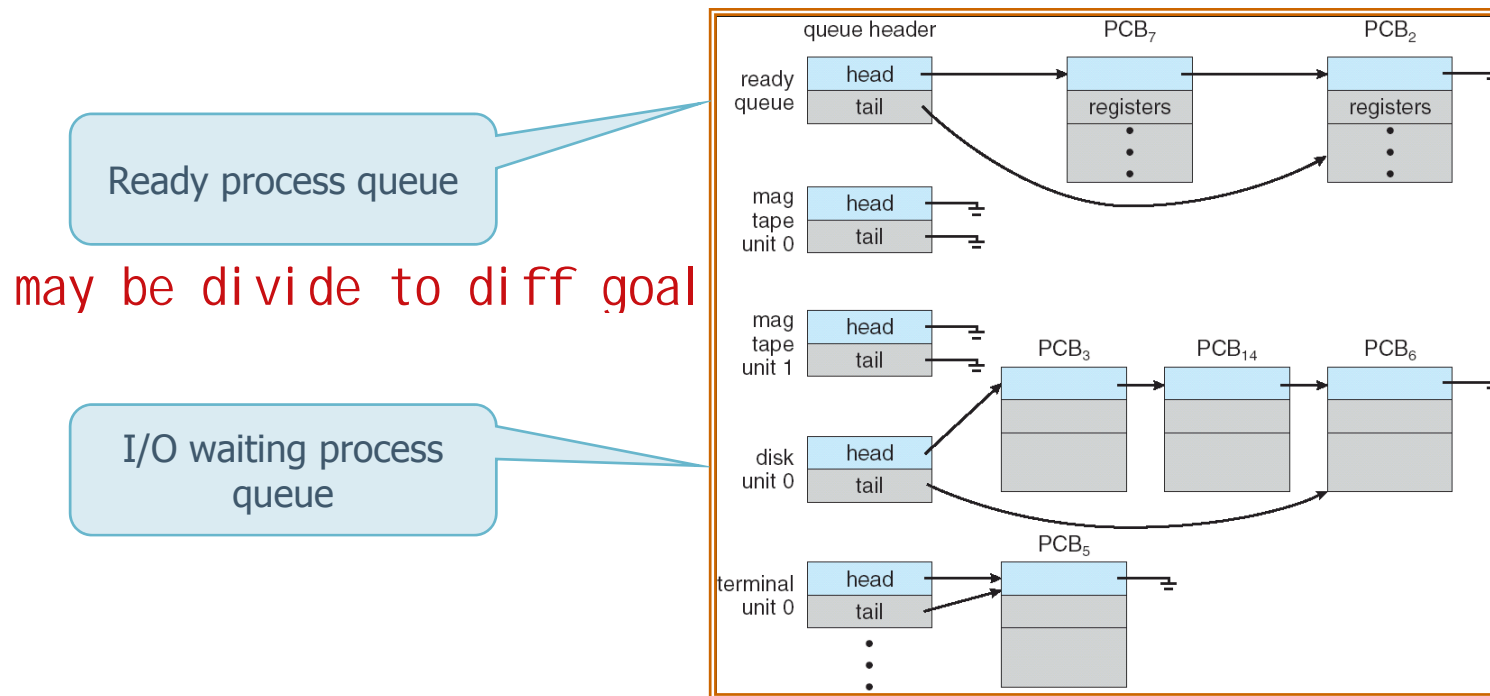
- Run less frequently
- Rescheduling time in the order of seconds/minutes
- Selects which process image can be loaded in main memory (swapper)
- Controls multiprogramming to avoid trashing
 - Too many processes that conflicts for the use of limited resources

high priority

Process scheduling

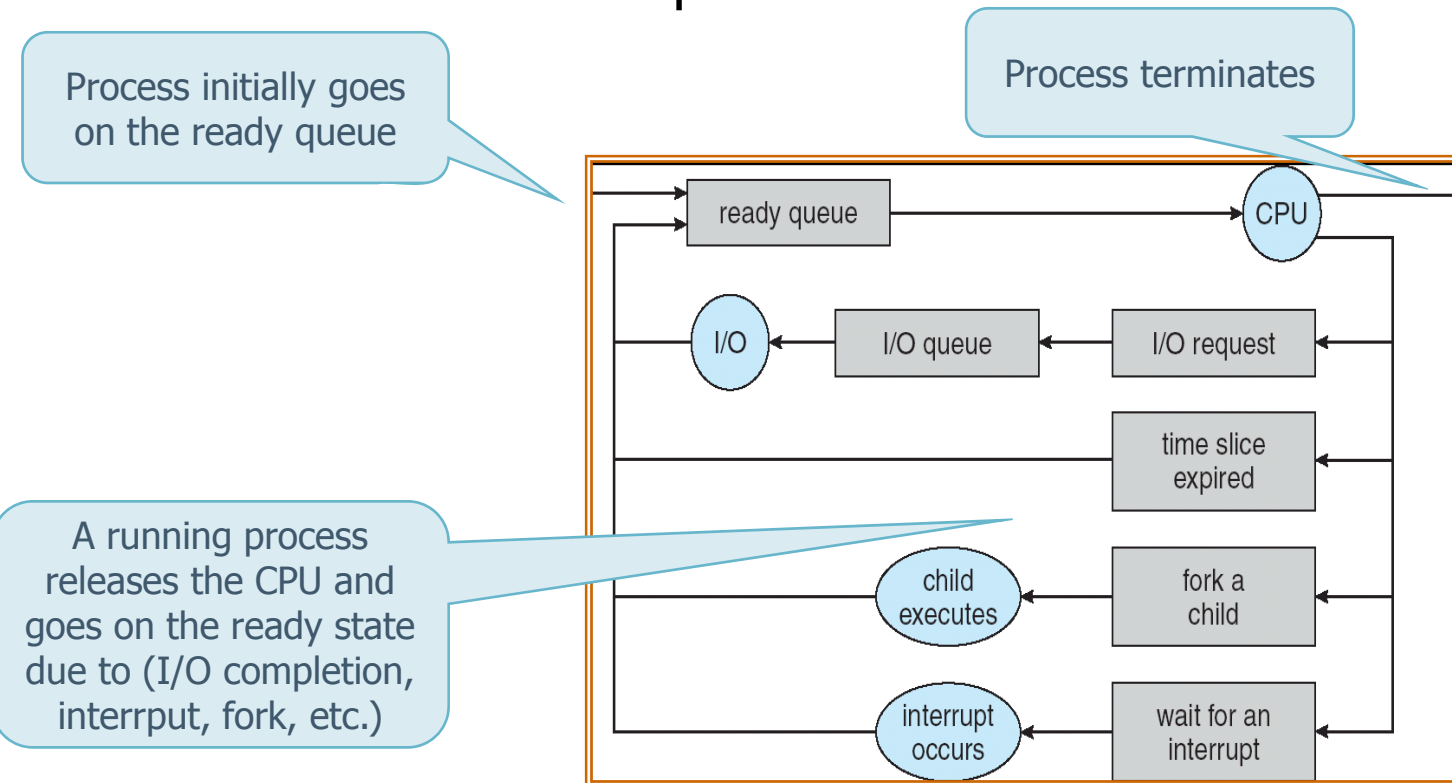
❖ A scheduler manages **process queues**

先进先出，正常排队



Queuing diagram

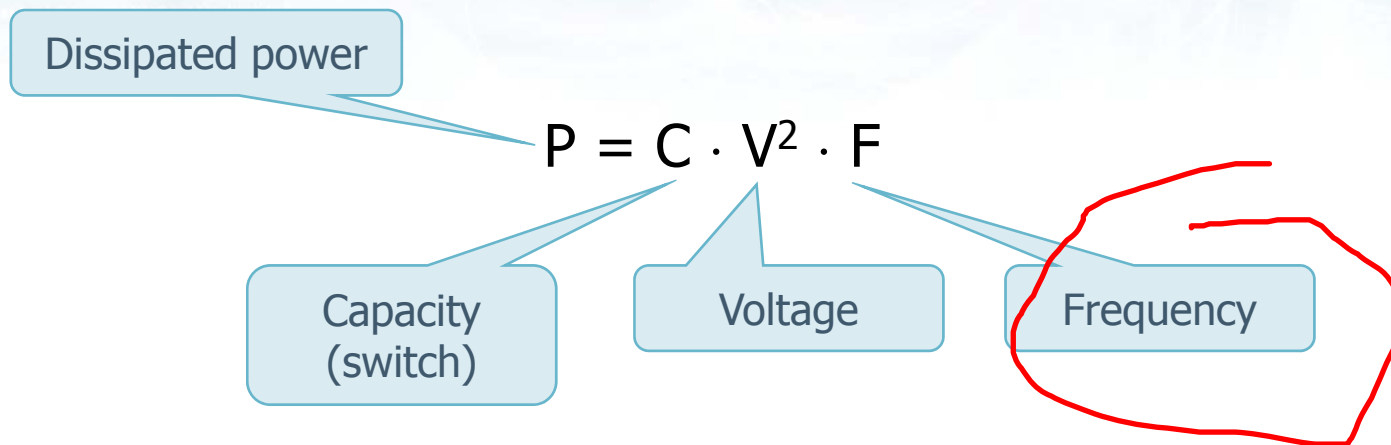
- ❖ The queuing diagram shows the possible process transitions from one queue to another one



Concurrency

- ❖ **Parallel computing** is a type of computation in which many tasks are carried out simultaneously.
 - Large problems can often be divided into smaller ones, which can then be solved at the same time
- ❖ **Concurrency**
 - Has been used since long time
 - Large development in the last years due to upper limits on the CPU frequency, and power consumption

Concurrency



- ❖ Due to the limits of frequency scaling
 - Parallel computing is now one of the main programming paradigms
 - Concurrent programming has introduced new challenges and pitfalls (bugs)

Concurrency

❖ Parallel computation levels

➤ Bit-level

- Word length determines the efficiency of an instruction (e.g., 8 bit versus 16 bit adder)

➤ Instruction-level

- Use of multi-stage pipelines for the execution of an instruction flow (e.g., fetch, decode, execute)

➤ Task-level

- Different computations are executed in parallel (e.g., sorting and matrix product)

Speed-up

- ❖ Concurrency aims at increasing the usage of the CPUs
 - Only multi-processor or multi-core systems allow obtaining real concurrency.
- ❖ Different levels of concurrency
 - Computer-cluster
 - Multi-processor
 - Multi-core

Speed-up

❖ **In theory** concurrency should allow linear speed-up

➤ Doubling the number of computation units, the execution time should be halved

❖ This behavior is obtained

- Rarely

- If a process is intrinsically sequential, augmenting the processing units does not change its execution time.

- For a limited number of processors/cores

- The speed-up curve grows linearly initially, but then goes towards an horizontal asymptote

difficult

Speed-up

❖ Amdahl law [1967]

- Small program segments intrinsically sequential limit the total speed-up that can be obtained

$$\lim_{P \rightarrow \infty} \frac{1}{\frac{1-\alpha}{P} + \alpha} = \frac{1}{\alpha}$$

#Processors

% sequential
run-time

$\alpha = 10\% = 0.1$

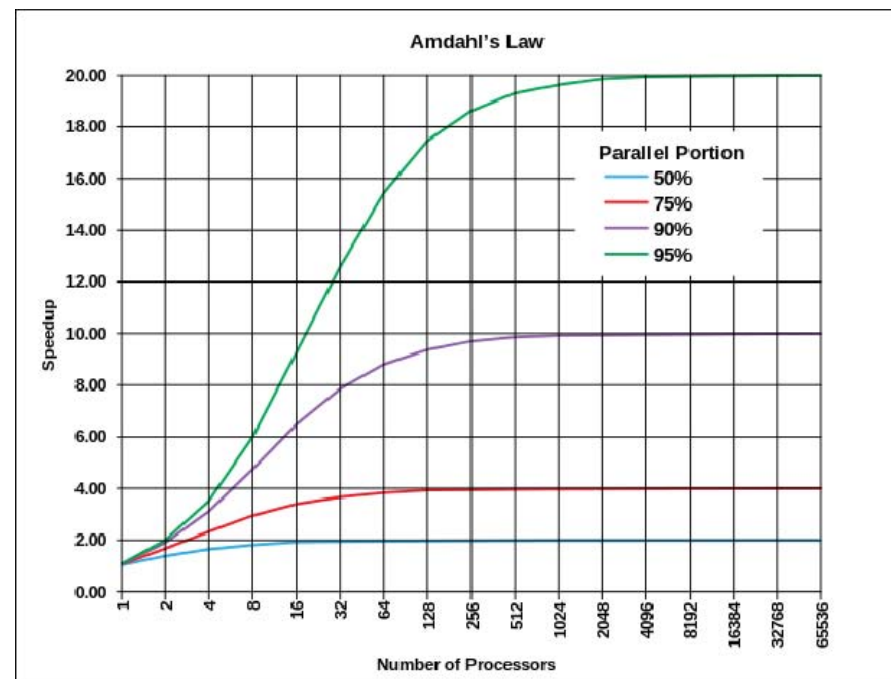
→

10x speed-up max

Speed-up

❖ Amdahl law [1967]

- Small program segments intrinsically sequential limit the total speed-up that can be obtained

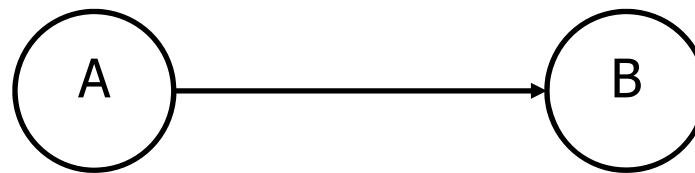


Parallel implementation

- ❖ An algorithm can be parallelized only respecting all its dependencies
 - Precedence constraints among operations (i.e., instructions, instruction blocks, processes)
 - A program cannot be executed faster than its slower sequence of operations (**critical path**)
- ❖ Precedence constraints can be represented by means of **precedence graphs**
 - Relation with **Control Flow Graph** and **Process generation trees**

Precedence graphs

- ❖ A precedence graph is an acyclic direct graph with
 - Nodes corresponding to instructions, instruction blocks, processes
 - Arcs represent precedence conditions
 - An arc from node A to node B means that B can be executed only when A is completed



Sequential and concurrent processes

❖ Sequential execution


➤ Actions are executed one **after** the other

- A new action begins only after the termination of the previous one

➤ Deterministic behavior

- Given the same input, the output produced is always the same, it does not depend on
 - The time of execution
 - The speed of execution
 - The number of active processes on the same system

Sequential
actions




Input,
Processing
Output

easily to debug

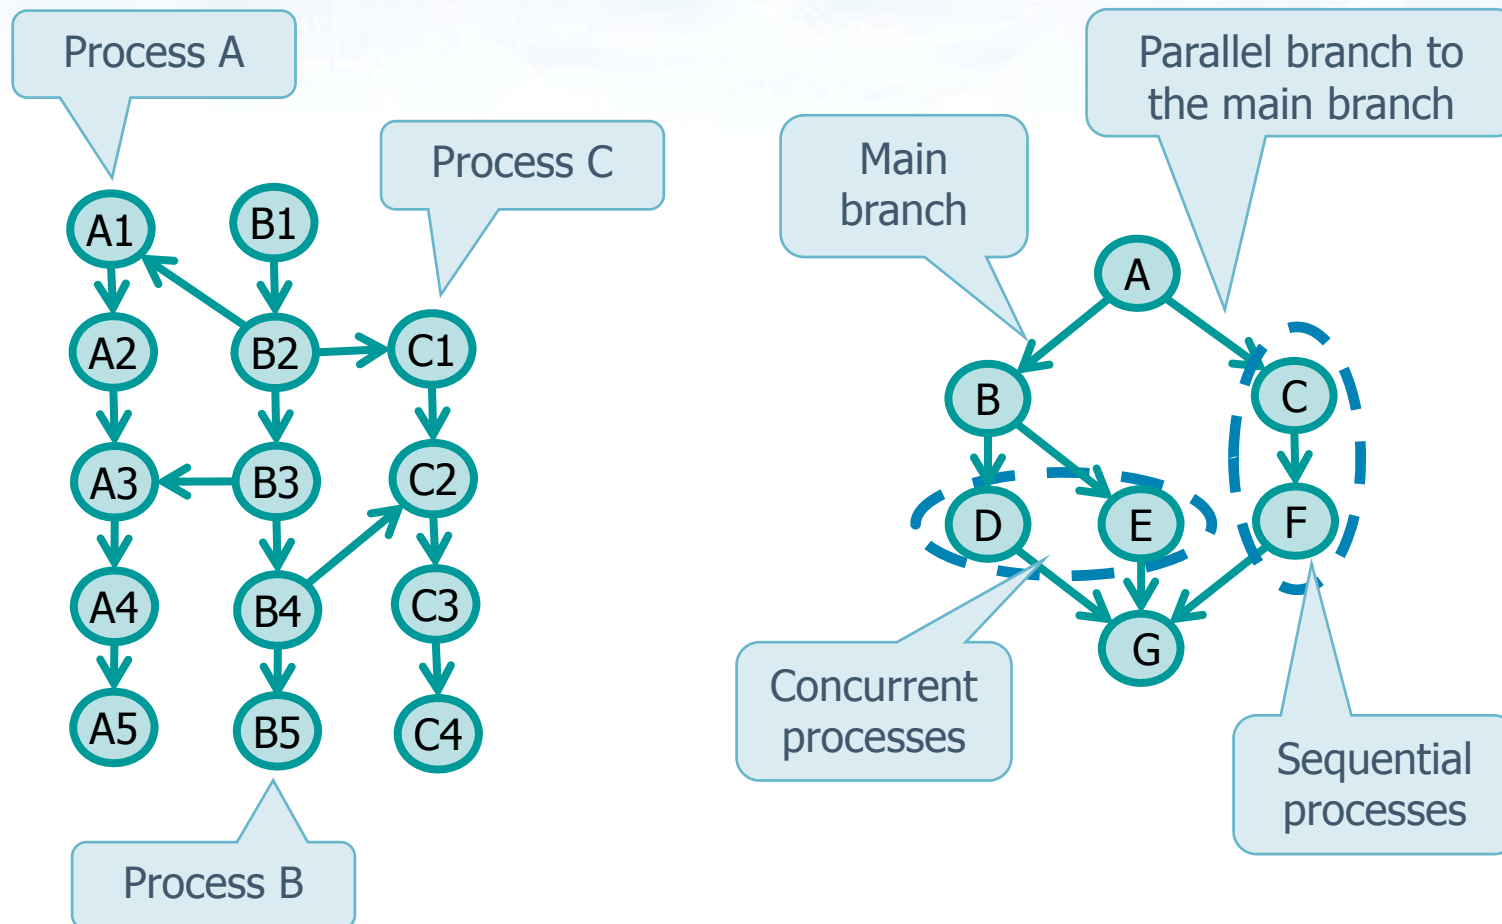
Sequential and concurrent processes

❖ Concurrent execution

- More than one action can be executed at the same time
 - There is not order relation
 - Non deterministic behavior
- Real concurrency
 - on multi-processor or multi-core systems
- Pseudo-concurrency
 - on mono-processor systems



Precedence graphs



Conditions for concurrency

hard to debug

- ❖ When two processes can be executed in concurrency?
- ❖ Given a process **P**, let's define
 - **R (P)**
 - Read set of P: the set of variables read by P
 - **W (P)**
 - Write set of P: the set of variables modified by P

Conditions for concurrency

❖ Bernstein conditions [1966]

- Two processes P_i and P_j can be executed in concurrency iff

eg, global var
several process

- $R(P_i) \cap W(P_j) = \emptyset$
- $W(P_i) \cap R(P_j) = \emptyset$
- $W(P_i) \cap W(P_j) = \emptyset$

R R is ok

- Otherwise time-dependent errors, or the programmer must impose regions of **Mutual Exclusion** among processes

Example

```
S1. a = x + y  
S2. b = z + 1  
S3. c = a - b  
S4. w = c + 1
```

❖ The sequential flow of instructions

➤ $S1 \rightarrow S2 \rightarrow S3 \rightarrow S4$

can be optimized because

- instruction 3 must be executed after 1 and 2
- instruction 4 must be executed after 3
- but, instructions 1 and 2 can be executed in parallel

Example

```
S1. a = x + y
S2. b = z + 1
S3. c = a - b
S4. w = c + 1
```

```
R(S1)={x, y}, W(S1)={a}
R(S2)={z},    W(S2)={b}
R(S3)={a, b}, W(S3)={c}
R(S4)={c},    W(S4)={w}
```

```
R(S1) ∩ W(S2) = 0
R(S2) ∩ W(S1) = 0
W(S1) ∩ W(S2) = 0
...
```

