

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
```

```
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE * f;
```

```
    for(i=0; i<MAXPAROLA; i++)
        freq[i]=0;
```

```
    if(argc != 2)
```

```
    {
        fprintf(stderr, "ERRORE: serve un parametro con il nome del file\n");
        exit(1);
    }
```

```
    f = fopen(argv[1], "r");
    if(f==NULL)
```

```
    {
        fprintf(stderr, "ERRORE: impossibile aprire il file %s\n", argv[1]);
        exit(1);
    }
```

```
    while( fgets( riga, MAXRIGA, f ) != NULL )
```

Processes

Signals

Stefano Quer and Pietro Laface

Dipartimento di Automatica e Informatica

Politecnico di Torino

Definition

- ❖ A **signal** is a software interrupt, i.e., an asynchronous notification sent, by the kernel or by another process, to a process to notify it of an event that occurred eg kernel

- ❖ Signals

- can be used as a limited form of inter-process communication
- allow notify asynchronous events such as the error conditions illustrated for exceptions

1. not other information
2. may be lost

Definition

- ❖ A **signal** is a software interrupt, i.e., an asynchronous notification sent, by the kernel or by another process, to a process to notify it of an event that occurred
- ❖ Signals
 - can be used as a limited form of inter-process communication
 - allow notify asynchronous events such as the error conditions illustrated for exceptions

Signals sent by the exception handlers

| Exception | Exception handler | Signal |
|----------------------|------------------------|---------|
| Divide error | divide_error() | SIGFPE |
| Debug | debug() | SIGTRAP |
| Breakpoint | int3() | SIGTRAP |
| Overflow | overflow() | SIGSEGV |
| Bounds check | bounds() | SIGSEGV |
| Invalid opcode | invalid_op() | SIGILL |
| Segment not present | segment_not_present() | SIGBUS |
| Stack segment fault | stack_segment() | SIGBUS |
| General protection | general_protection() | SIGSEGV |
| Page Fault | page_fault() | SIGSEGV |
| Intel-reserved | None | None |
| Floating-point error | coprocessor_error() | SIGFPE |

Terminal signals

- ❖ Typing some key combinations at the controlling terminal of a process causes the system to send it a signal:
- ❖ Ctrl-C sends an **SIGINT** signal
 - by default, this causes the process to terminate.
- ❖ Ctrl-Z sends a terminal stop SIGTSTP signal
 - by default, this causes the process to suspend execution.
- ❖ Ctrl-\ sends a **SIGQUIT** signal;
 - by default, this causes the process to terminate and dump core.

Main signals

| Name | Description |
|----------------------------------|---|
| SIGABRT | Process abort, generated by system call abort |
| SIGALRM | Alarm clock, generated by system call alarm |
| SIGFPE | Floating-Point exception |
| SIGILL | Illegal instruction |
| SIGKILL | Kill (non maskable) |
| SIGPIPE | Write on a pipe with no reader |
| SIGSEGV | Invalid memory segment access |
| SIGCHLD | Child process stopped or exited |
| SIGUSR1 SIGUSR2 | User-defined signal 1/2 for sync |

You can display the complete list of signals using the shell command `kill -l`

Signal management

- ❖ Signal management goes through three phases:
signal generation, signal delivery, reaction to a signal

interrupt manage

- Signal generation

- When the kernel or a process causes an event that generate a signal

`exit(0)`

- Signal delivery

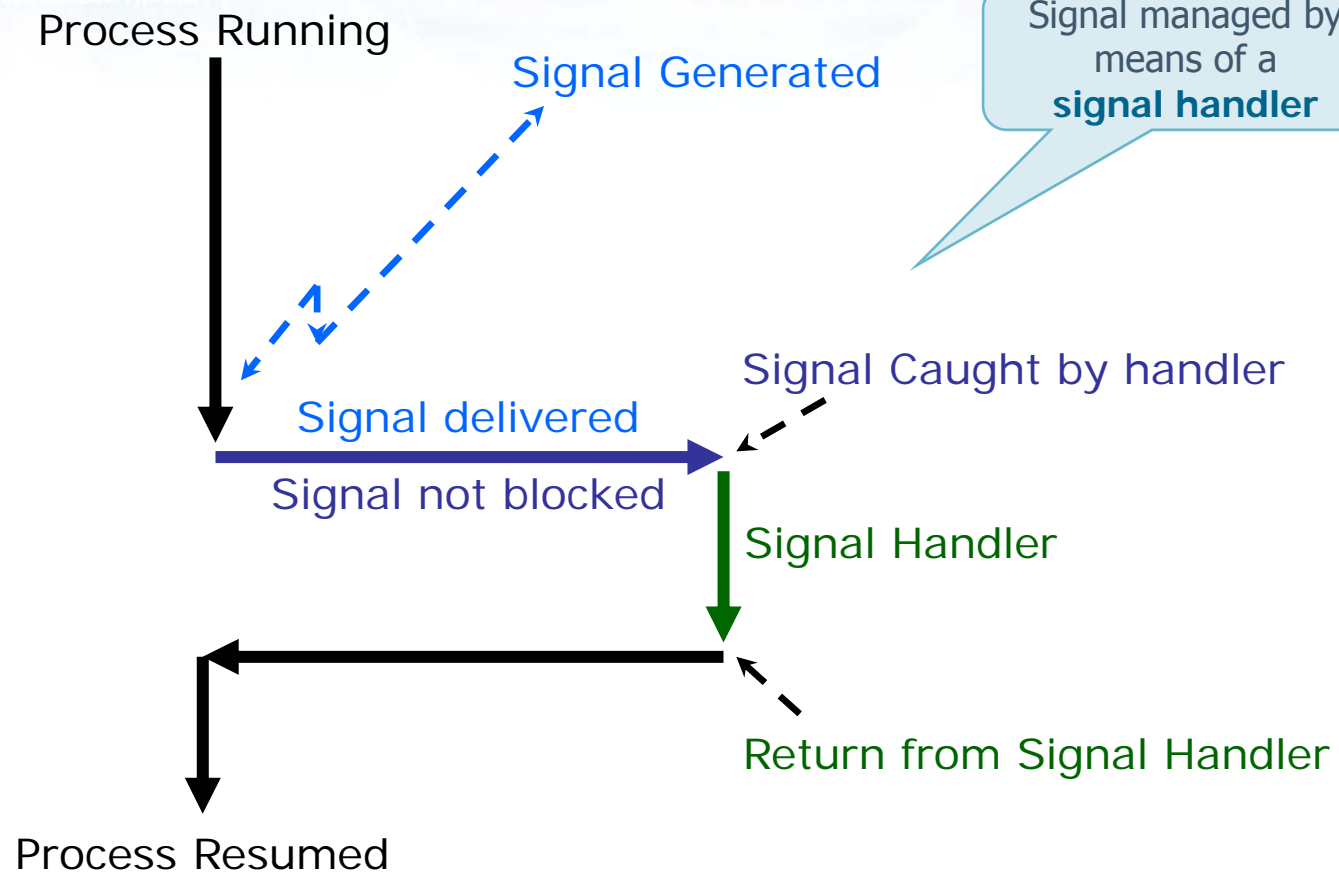
- A not yet delivered signal remains pending
- At signal delivery a process executes the actions related to that signal
- The lifetime of a signal is from its generation to its delivery

Signal management

➤ Reaction to a signal

- To properly react to the asynchronous arrival of a given type of signal, a process must inform the kernel about the action that it will perform when it will receive a signal of that type
- A process may
 - **Accept** the **default** behavior (be terminated)
 - Declare to the kernel that it wants to **ignore** the signals of that type
 - Declare to the kernel that it wants to **catch** and manage the signals of that type by means of a signal handler function (similarly to the interrupt management)

Signal management



Signal management

- ❖ Precondition to properly handle a received signal for a process is to **declare** to the kernel if a signal of a given type will be ignored or caught.
- ❖ This is done using the system call **signal**
 - Which **instantiates** a signal handler

signal does not send a signal !!

send by kill ()

signal() system call

```
#include <signal.h>

void (*signal (int sig,
               void (*func)(int)))(int);
```

❖ Arguments

tell what sig to do not send sig

- **sig** indicates the type of signal to be caught
 - SIGALRM, SIGUSR1, etc.
- **func** specifies the address (i.e., pointer) to the function that will be executed when a signal of that type is received by the process
 - This function has a single argument of `int` type, which indicates the type of signal that will be handled

signal() system call

```
#include <signal.h>

void (*signal (int sig,
                void (*func)(int)))(int);
```

❖ Return value

- the previous value of the signal handler, i.e., the previous signal handler function
- **SIG_ERR** in case of error, **errno** is set to indicate the cause

Signal generation

- ❖ The kernel generates signals
 - SIGCHLD, SIGFPE, etc.
- ❖ A process can (ask the kernel to) generate a signal by means of the system call

- **kill** (and **raise**)

`self kill`

kill is misleading, does not kill a process, just send to it a signal

- **alarm**

- Ask the kernel to receive a SIGALRM after a specified amount of time

`eg sleep()`

Waiting for a signal

- ❖ A process can wait for a signal by means of the system call
 - **pause**, and any other blocking system call
 - Suspend the process until any signal is received
 - **sleep**
 - Suspend the process for a specified amount of time (waits for signal SIGALRM)

Reaction to a signal

- ❖ **signal** system call allows setting three different reactions to the delivery of a signal
 - Accept the default behavior
 - signal (SIGname, **SIG_DFL**)
 - Where **SIG_DFL** is defined in `signal.h`
 - `#define SIG_DFL ((void (*)(void)) 0)`
 - Every signal has its own default behavior, defined by the system
 - Most of the default reactions is **process termination**

Reaction to a signal

➤ Ignore signal delivery

- signal (SIGname, **SIG_IGN**)
- Where **SIG_IGN** is defined in `signal.h`
 - `#define SIG_DFL ((void (*)(void)) 1)`
- Some signals cannot be ignored
 - `SIGKILL` and `SIGSTOP` cannot be ignored because the kernel and the superuser would not have the possibility to control all processes
 - Ignoring an illegal memory access, signaled by `SIGSEGV`, would produce an undefined process behavior

eg when using &
send it to kernel to
prevent zombie process

Reaction to a signal

➤ Catch the signal

- signal (SIGname, signalHandlerFunction)
- where
 - **SIGname** indicates the signal type
 - **signalHandlerFunction** is the user defined signal handler function
- The signal handler
 - Is executed when the signal is delivered,
 - When it returns, the process continues with the next instruction, as it happens for interrupts

A signal handler function must be defined for every signal type that must be caught

Example 1

```
#include <signal.h>
#include <stdio.h>
#include <unistd.h>

void manager (int sig) {
    printf ("Received signal %d\n", sig);
    // signal (SIGINT, manager);
    return;
}

int main() {
    signal (SIGINT, manager);
    while (1) {
        printf ("main: Hello!\n");
        sleep (1);
    }
}
```

Signal handler for
signal SIGINT

Declares the signal
handler

Example 2

```
...
void manager (int sig) {
    if (sig==SIGUSR1)
        printf ("Received SIGUSR1\n");
    else if (sig==SIGUSR2)
        printf ("Received SIGUSR2\n");
    else printf ("Received %d\n", sig);
    return;
}
...
int main () {
    ...
    signal (SIGUSR1, manager);
    signal (SIGUSR2, manager);
    ...
}
```

Same signal handler
for more than one
signal type

Both signal types
must be declared

Example 3

Standard behavior of **wait**

```
// signal (SIGCHLD, SIG_IGN);

for (i=0; i<3; i++) {
    if (fork() == 0) {
        // child
        sleep (1);
        printf ("i=%d PID=%d\n", i, getpid());
        exit (i);
    }
}
sleep (5);
for (i=0; i<3; i++) {
    ret = wait (&code);
    printf ("Wait: ret=%d code=%x\n", ret, code);
}
```

i=2 PID=3057
i=1 PID=3056
i=0 PID=3055

Wait: ret = 3055 code = 0
Wait: ret = 3056 code = 100
Wait: ret = 3057 code = 200

Example 3

Altering the behavior of
wait

```
signal (SIGCHLD, SIG_IGN);
```

Ignore SIGCHLD, sent
by the kernel to the
parent at the exit of a
child

```
for (i=0; i<3; i++) {  
    if (fork() == 0) {  
        // child  
        sleep (1);  
        printf ("i=%d PID=%d\n", i, getpid());  
        exit (i);  
    }  
}  
sleep (5);  
for (i=0; i<3; i++) {  
    ret = wait (&code);  
    printf ("Wait: ret=%d code=%x\n", ret, code);  
}
```

No wait:
Wait: ret = -1 code = 7FFF
Wait: ret = -1 code = 7FFF
Wait: ret = -1 code = 7FFF

kill system call

```
#include <signal.h>

int kill (pid_t pid, int sig);
```

- ❖ Send signal (**sig**) to a process or to a group of processes (**pid**)
 - A **user** process can send signals only to processes having the same UID
 - The **superuser** can send signal to any process

kill system call

```
#include <signal.h>

int kill (pid_t pid, int sig);
```

❖ Arguments

| If pid is | Send sig |
|-----------|---|
| >0 | To process with PID equal to <code>pid</code> |
| ==0 | To all processes with GID equal to its GID |
| <0 | To all processes with GID equal to the absolute value of <code>pid</code> |
| == -1 | To all processes |

"All process" excludes
a set of system
processes

kill system call

```
#include <signal.h>

int kill (pid_t pid, int sig);
```

❖ Return value

- 0 on success
- -1 on error

raise system call

```
#include <signal.h>

int raise (int sig);
```

- ❖ The **raise** system call allows a process to send a signal to itself
 - `raise (sig)` is equivalent to
 - `kill (getpid(), sig)`

pause system call

```
#include <unistd.h>

int pause (void);
```

- ❖ Suspends the calling process until a signal is delivered
- ❖ Returns after the completion of the signal handler
 - returns -1, and errno is set to EINTR

kill -9

alarm system call

```
#include <unistd.h>

unsigned int alarm (unsigned int seconds);
```

- ❖ Ask the kernel so send a **SIGALRM** to the calling process in **seconds** seconds
 - The default action for SIGALRM is process termination
 - A call to **alarm(seconds)** before the expiration of a previous **alarm** reschedules the request to the kernel
 - **alarm(0)** cancels any pending alarm

alarm system call

```
#include <unistd.h>

unsigned int alarm (unsigned int seconds);
```

❖ Return value

- the number of seconds remaining until the delivery of a previously scheduled alarm
- zero if there was no previously scheduled alarm.

alarm system call

```
#include <unistd.h>

unsigned int alarm (unsigned int seconds);
```

❖ Warning

- The signal is generated by the kernel
 - It is possible that the process get the CPU control after some time, depending on the scheduler decisions
- System calls **sleep** and **alarm** uses the same kernel timer

dont use at the same time

Example

- ❖ Implement system call **sleep** using system calls **alarm** and **pause**

```
include <signal.h>
#include <unistd.h>

static void sig_alarm(int signo) {return;}

unsigned int sleep1(unsigned int nsecs)
{
    if (signal(SIGALRM, sig_alarm) == SIG_ERR)
        return (nsecs);
    alarm (nsecs);
    pause ();
    return (alarm(0));
}
```

Returns 0, or the remaining time before the delivery if **pause** returns because another signal has been received

Example

- ❖ Implement system call **alarm** using system calls **fork**, **signal**, **kill** and **pause**

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void myAlarm (int sig) {
    printf ("Alarm\n");
}
```

Example

```
int main (void) {
    pid_t pid;
    (void) signal (SIGALRM, myAlarm);
    pid = fork();
    switch (pid) {
        case -1: /* error */
            printf ("fork failed");
            exit (1);
        case 0: /* child */
            sleep(5);
            kill (getppid(), SIGALRM);
            exit(0);
    }
    /* parent */
    pause ();
    exit (0);
}
```

The child waits
and sends
SIGALRM

The parent pauses, and continues
only when it receives the SIGALRM
sent by the child

Signal limitations

- ❖ Signals do not convey any information
- ❖ The **memory** of the pending signals is **limited**
 - Max one signal pending per type
 - Forthcoming signals of the same type are lost
 - Signals can be ignored
- ❖ Signal interrupt functions that must be **reentrant**
- ❖ Produce **race conditions**
- ❖ Some limitations are avoided in POSIX.4

Limited memory

```
...
static void sigUshr1 (int);
static void sigUshr2 (int);

static void
sigUshr1 (int signo) {
    if (signo == SIGUSR1)
        printf("Received SIGUSR1\n");
    else
        printf("Received wrong SIGNAL\n");

    fprintf (stdout, "sigUshr1 sleeping ...\n");
    sleep (5);
    fprintf (stdout, "... sigUshr1 end sleeping.\n");
    return;
}
```

Limited memory

```
static void
sigUshr2 (int signo) {
    if (signo == SIGUSR2)
        printf("Received SIGUSR2\n");
    else
        printf("Received wrong SIGNAL\n");

    fprintf (stdout, "sigUshr2 sleeping ...\n");
    sleep (5);
    fprintf (stdout, "... sigUshr2 end sleeping.\n");

    return;
}
```

Limited memory

```
int
main (void) {
    if (signal(SIGUSR1, sigUusr1) == SIG_ERR) {
        fprintf (stderr, "Signal Handler Error.\n");
        return (1);
    }
    if (signal(SIGUSR2, sigUusr2) == SIG_ERR) {
        fprintf (stderr, "Signal Handler Error.\n");
        return (1);
    }
    while (1) {
        fprintf (stdout, "Before pause.\n");
        pause ();
        fprintf (stdout, "After pause.\n");
    }
    return (0);
}
```

The main iterates waiting signals from shell

Limited memory

Shell commands

```
> ./pgrm &  
[3] 2636  
> Before pause.  
> kill -USR1 2636  
> Received SIGUSR1  
sigUsr1 sleeping ...  
... sigUsr1 end sleeping.  
After pause.  
Before pause.  
> kill -USR2 2636  
> Received SIGUSR2  
sigUsr2 sleeping ...  
... sigUsr2 end sleeping.  
After pause.  
Before pause.
```

Correctly received
SIGUSR1

Correctly received
SIGUSR2

Observation:
shell command **kill** sends a signal to
a process with a specified PID

Limited memory

```
> kill -USR1 2636 ; kill -USR2 2636
> Received SIGUSR2
sigUusr2 sleeping ...
... sigUusr2 end sleeping.
Received SIGUSR1
sigUusr1 sleeping ...
... sigUusr1 end sleeping.
After pause.
Before pause.
```

Two signals sent in sequence

Both are received

The deliver order of the two signal cannot be predicted

Limited memory

not received

```
> kill -USR1 2636 ; kill -USR2 2636 ; kill -USR1 2636
> Received SIGUSR1
sigUusr1 sleeping ...
... sigUusr1 end sleeping.
Received SIGUSR2
sigUusr2 sleeping ...
... sigUusr2 end sleeping.
After pause.
Before pause.

> kill -9 2636
[3]+  Killed  ./pgrm
```

Three signals sent in sequence: two of them are SIGUSR1

A SIGUSR1 is lost

-9 = SIGKILL = Kill
Kill a process

Reentrant functions

- ❖ The kernel **knows** where a signal handler returns, but
- ❖ The signal handler **does not know** where it was called, i.e., the control flow was interrupted by the signal

Reentrant functions: Examples

- ❖ Suppose a **malloc** is interrupted, and the signal handler calls another malloc
 - Function malloc manages the list of the free memory regions, which could be corrupted
- ❖ Suppose that the execution of a function that uses a **static variable** is interrupted, but is then called by the signal handler
 - The static variable could be used to store a new value, i.e., it does not remain the same it was before the signal was delivered

Reentrant functions: Conclusions

- ❖ The Single UNIX Specification defines the reentrant functions, which can be interrupted without problems
 - read, write, sleep, wait, etc.
- ❖ Most of the I/O standard C are not reentrant
 - printf, scanf, etc.
 - They use static variables or global variables, thus must be used carefully and being aware of possible problems

Race conditions

❖ Race condition

- The result of more concurrent processes working on common data depends on the execution order of the processes instructions
- ❖ Using signals increases the probability of race conditions.

Race conditions example

- ❖ Suppose a process decides to suspend itself for a given number of seconds
- ❖ The signal could be delivered before the execution of pause due to a contest switching and scheduling decisions.

```
static void  
myHandler (int signo) {  
    ...  
}  
...  
signal (SIGALRM, myHandler)  
alarm (nSec);  
pause ();
```

Signal **SIGALRM** can be delivered before **pause**


pause blocks the process forever because the signal has been lost

Race conditions example

- ❖ Suppose two processes P_1 and P_2 decide to synchronize by means of signals
 - If P_1 (P_2) signal is delivered before P_2 (P_1) executes **pause**
 - Process P_2 (P_1) blocks forever waiting a signal


```
P1
while (1) {
    ...
    kill (pidP2, SIG...);
    pause ();
}
```

```
P2
while (1) {
    pause ();
    ...
    kill (pidP1, SIG...);
}
```



Exercise

- ❖ Despite their shortcomings, signals can provide a rough synchronization mechanism
- ❖ **Ignoring the race conditions** (and using `fork`, `wait`, `signal`, `kill`, and `pause`) implement this precedence graph



Solution

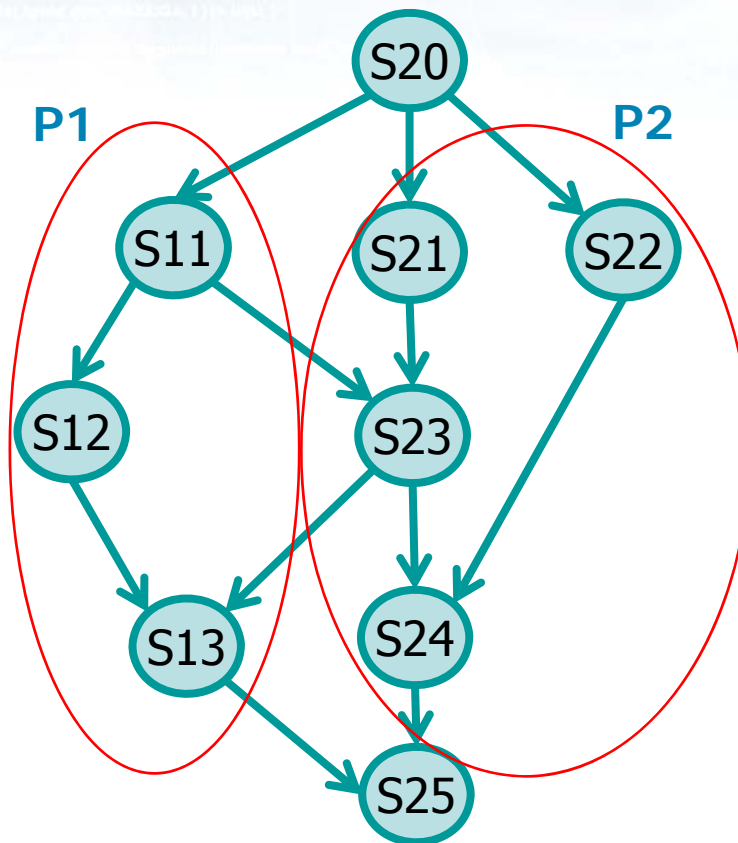
```
static void
sigUshr ( int signo) {
    if (signo== SIGUSR1)
        printf ("SIGUSR1\n");
    else if (signo==SIGUSR2)
        printf ("SIGUSR2\n");
    else
        printf ("Signal %d\n", signo);
    return;
}
```

Solution

```
int main (void) {
    pid_t pid;

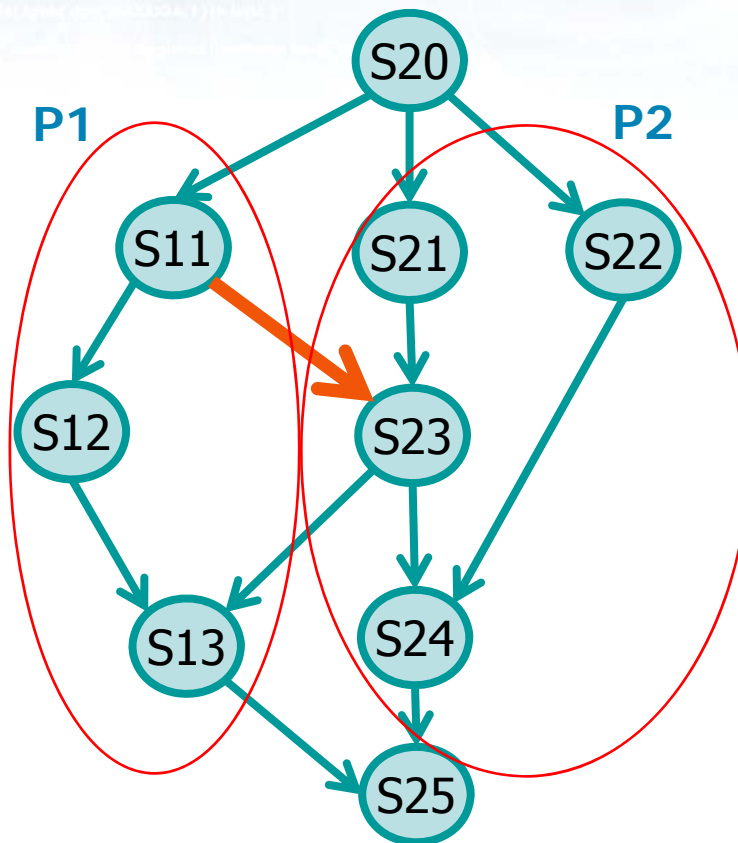
    if (signal(SIGUSR1, sigUshr) == SIG_ERR) {
        printf ("Signal Handler Error.\n");
        return (1);
    }
    if (signal(SIGUSR2, sigUshr) == SIG_ERR) {
        printf ("Signal Handler Error.\n");
        return (1);
    }
}
```


Solution



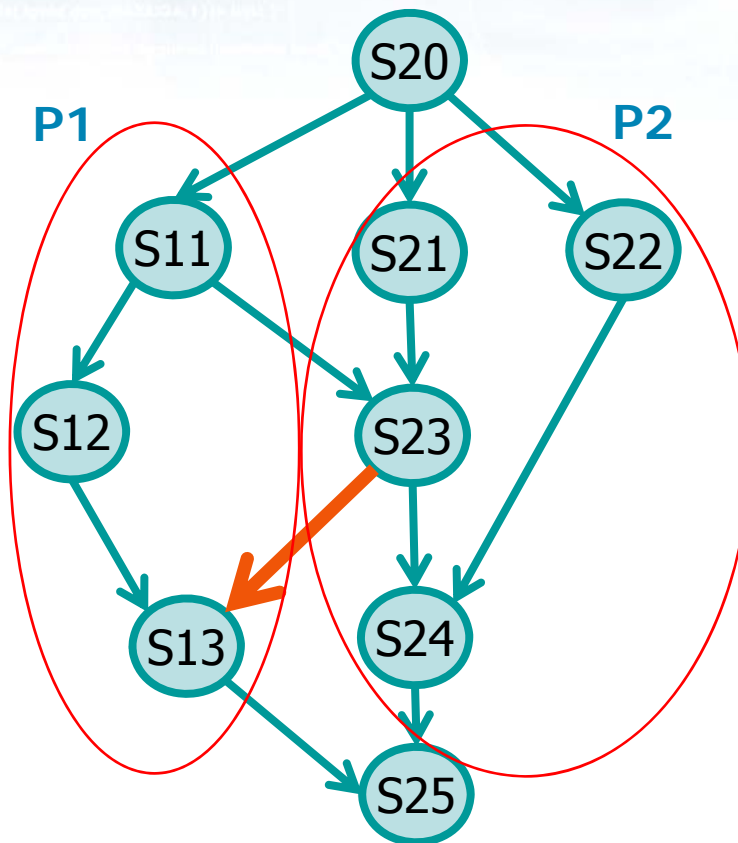
```
printf ("S20\n");
pid = fork ();
if (pid > (pid_t) 0) {
    P1 (pid);
    wait ((int *) 0);
} else {
    P2 ();
    exit (0);
}
printf ("S25\n");
return (0);
}
```

Solution



```
void P1 (  
    pid_t cpid  
) {  
    printf ("S11\n");  
    sleep (1);  // !?  
    kill (cpid, SIGUSR1);  
    printf ("S12\n");  
    pause ();  
    printf ("S13\n");  
  
    return;  
}
```

Solution



```
void P2 ( ) {  
    if (fork ( ) > 0) {  
        printf ("S21\n");  
        pause ( );  
        printf ("S23\n");  
        kill (getppid ( ),  
              SIGUSR2);  
        wait ((int *) 0);  
    } else {  
        printf ("S22\n");  
        exit (0);  
    }  
    printf ("S24\n");  
    return;  
}
```